

In [27]:

```
# import packages
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

import string
import random

from sklearn.preprocessing import StandardScaler
from itertools import product
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import RepeatedStratifiedKFold

from sklearn import tree
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import KFold
from sklearn.metrics import classification_report
from sklearn.metrics import accuracy_score
from sklearn.preprocessing import OneHotEncoder, LabelEncoder
from sklearn.metrics import roc_curve, auc
from sklearn.tree import DecisionTreeRegressor, DecisionTreeClassifier
from sklearn.model_selection import train_test_split

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import GradientBoostingClassifier

import sklearn.metrics as metrics

from matplotlib import pyplot
from matplotlib.pyplot import figure

from sklearn.model_selection import RandomizedSearchCV

from sklearn.model_selection import GridSearchCV
from xgboost import XGBClassifier

from scipy.stats import loguniform
from sklearn import metrics

from sklearn.ensemble import AdaBoostClassifier
from sklearn.feature_selection import RFE

import time
import datetime

from tqdm import tqdm

from imblearn.ensemble import EasyEnsembleClassifier

from imblearn.ensemble import RUSBoostClassifier

from imblearn.datasets import make_imbalance

from sklearn.model_selection import cross_validate

from sklearn.dummy import DummyClassifier

from imblearn.ensemble import BalancedRandomForestClassifier

import warnings
from sklearn.exceptions import ConvergenceWarning

from imblearn.ensemble import BalancedBaggingClassifier
```

```

from imblearn.over_sampling import SMOTE

from sklearn.ensemble import HistGradientBoostingClassifier
def fxn():
    warnings.warn("deprecated", DeprecationWarning)
    warnings.warn("future", FutureWarning)
    warnings.warn("convergence", ConvergenceWarning)

with warnings.catch_warnings():
    warnings.simplefilter("ignore")
    fxn()

from sklearn import *

```

## Initialization

```

In [2]: # set up google cloud to store the future records of hypertune
import gspread
sa = gspread.service_account(filename="kaggle-hypertune-records-1d18a6abfb85.json")

```

```

In [98]: # import data
train_data = pd.read_csv('train.csv')
test_data = pd.read_csv('test.csv')
sample_submission_data = pd.read_csv('sample_submission.csv')

```

```

In [99]: # basic information
train_data.info()

# describe the data
train_data.head()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1340 entries, 0 to 1339
Data columns (total 35 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   EmployeeID                           1340 non-null   int64
 1   Age                                   1340 non-null   int64
 2   Attrition                           1340 non-null   object
 3   BusinessTravel                       1340 non-null   object
 4   DailyRate                           1340 non-null   int64
 5   Department                           1340 non-null   object
 6   DistanceFromHome                    1340 non-null   int64
 7   Education                            1340 non-null   int64
 8   EducationField                       1340 non-null   object
 9   EmployeeCount                       1340 non-null   int64
10  EnvironmentSatisfaction              1340 non-null   int64
11  Gender                               1340 non-null   object
12  HourlyRate                           1340 non-null   int64
13  JobInvolvement                       1340 non-null   int64
14  JobLevel                             1340 non-null   int64
15  JobRole                              1340 non-null   object
16  JobSatisfaction                      1340 non-null   int64
17  MaritalStatus                       1340 non-null   object
18  MonthlyIncome                       1340 non-null   int64
19  MonthlyRate                         1340 non-null   int64
20  NumCompaniesWorked                  1340 non-null   int64
21  Over18                              1340 non-null   object
22  OverTime                            1340 non-null   object

```

```
23 PercentSalaryHike      1340 non-null int64
24 PerformanceRating      1340 non-null int64
25 RelationshipSatisfaction 1340 non-null int64
26 StandardHours          1340 non-null int64
27 Shift                  1340 non-null int64
28 TotalWorkingYears      1340 non-null int64
29 TrainingTimesLastYear  1340 non-null int64
30 WorkLifeBalance        1340 non-null int64
31 YearsAtCompany         1340 non-null int64
32 YearsInCurrentRole     1340 non-null int64
33 YearsSinceLastPromotion 1340 non-null int64
34 YearsWithCurrManager   1340 non-null int64
```

dtypes: int64(26), object(9)

memory usage: 366.5+ KB

```
Out[99]:
```

	EmployeeID	Age	Attrition	BusinessTravel	DailyRate	Department	DistanceFromHome	Education	Educ
0	1317087	40	No	Travel_Rarely	1398	Cardiology	2	4	Li
1	1548175	40	No	Travel_Rarely	1300	Maternity	24	2	
2	1215433	25	No	Travel_Rarely	622	Cardiology	13	1	
3	1375351	33	No	Travel_Rarely	922	Maternity	1	5	
4	1028734	39	No	Travel_Frequently	505	Maternity	2	4	

5 rows x 35 columns

```
In [100...  
# check for missing data  
train_data.isnull().sum()  
test_data.isnull().sum()  
# there is no missing data
```

```
Out[100...  
EmployeeID      0  
Age             0  
BusinessTravel  0  
DailyRate       0  
Department      0  
DistanceFromHome 0  
Education       0  
EducationField  0  
EmployeeCount   0  
EnvironmentSatisfaction 0  
Gender          0  
HourlyRate      0  
JobInvolvement  0  
JobLevel        0  
JobRole         0  
JobSatisfaction 0  
MaritalStatus   0  
MonthlyIncome   0  
MonthlyRate     0  
NumCompaniesWorked 0  
Over18          0  
OverTime        0  
PercentSalaryHike 0  
PerformanceRating 0  
RelationshipSatisfaction 0  
StandardHours   0  
Shift           0  
TotalWorkingYears 0  
TrainingTimesLastYear 0  
WorkLifeBalance 0
```

```
YearsAtCompany      0
YearsInCurrentRole   0
YearsSinceLastPromotion 0
YearsWithCurrManager 0
dtype: int64
```

```
In [6]: # check for duplicated data
train_data.duplicated().sum()
train_data.duplicated().sum()
```

Out[6]: 0

```
In [82]: # feature engineering
# Categorical -> Numerical

def quantifyCategoricalTest(dataset):
    dataset = pd.get_dummies(dataset, columns=["Department", "EducationField", "JobRole", "MaritalStatus", "BusinessTravel"])
    dataset['BusinessTravel'].replace(['Non-Travel', 'Travel_Rarely', 'Travel_Frequently'], [0, 1, 2], inplace=True)
    dataset['Gender'].replace(['Female', 'Male'], [0, 1], inplace=True)
    dataset['Over18'].replace(['N', 'Y'], [0, 1], inplace=True)
    dataset['OverTime'].replace(['No', 'Yes'], [0, 1], inplace=True)

    return dataset

def quantifyCategoricalTrain(dataset):
    dataset = quantifyCategoricalTest(dataset)
    dataset['Attrition'].replace(['No', 'Yes'], [0, 1], inplace=True)

    return dataset
```

```
In [101]: train_data = quantifyCategoricalTrain(train_data)
test_data = quantifyCategoricalTest(test_data)
```

```
In [102]: # drop data
train_data = train_data.drop(['EmployeeID', 'EmployeeCount', 'Over18', 'StandardHours'], axis=1)
test_data = test_data.drop(['EmployeeID', 'EmployeeCount', 'Over18', 'StandardHours'], axis=1)
```

```
In [10]: train_data.head()
```

Out[10]:

	Age	Attrition	BusinessTravel	DailyRate	DistanceFromHome	Education	EnvironmentSatisfaction	Gender
0	40	0	1	1398	2	4	3	0
1	40	0	1	1300	24	2	1	1
2	25	0	1	622	13	1	2	1
3	33	0	1	922	1	5	1	0
4	39	0	2	505	2	4	3	0

5 rows x 40 columns

```
In [11]: train_data.describe()
```

Out[11]:

	Age	Attrition	BusinessTravel	DailyRate	DistanceFromHome	Education	Environment
--	-----	-----------	----------------	-----------	------------------	-----------	-------------

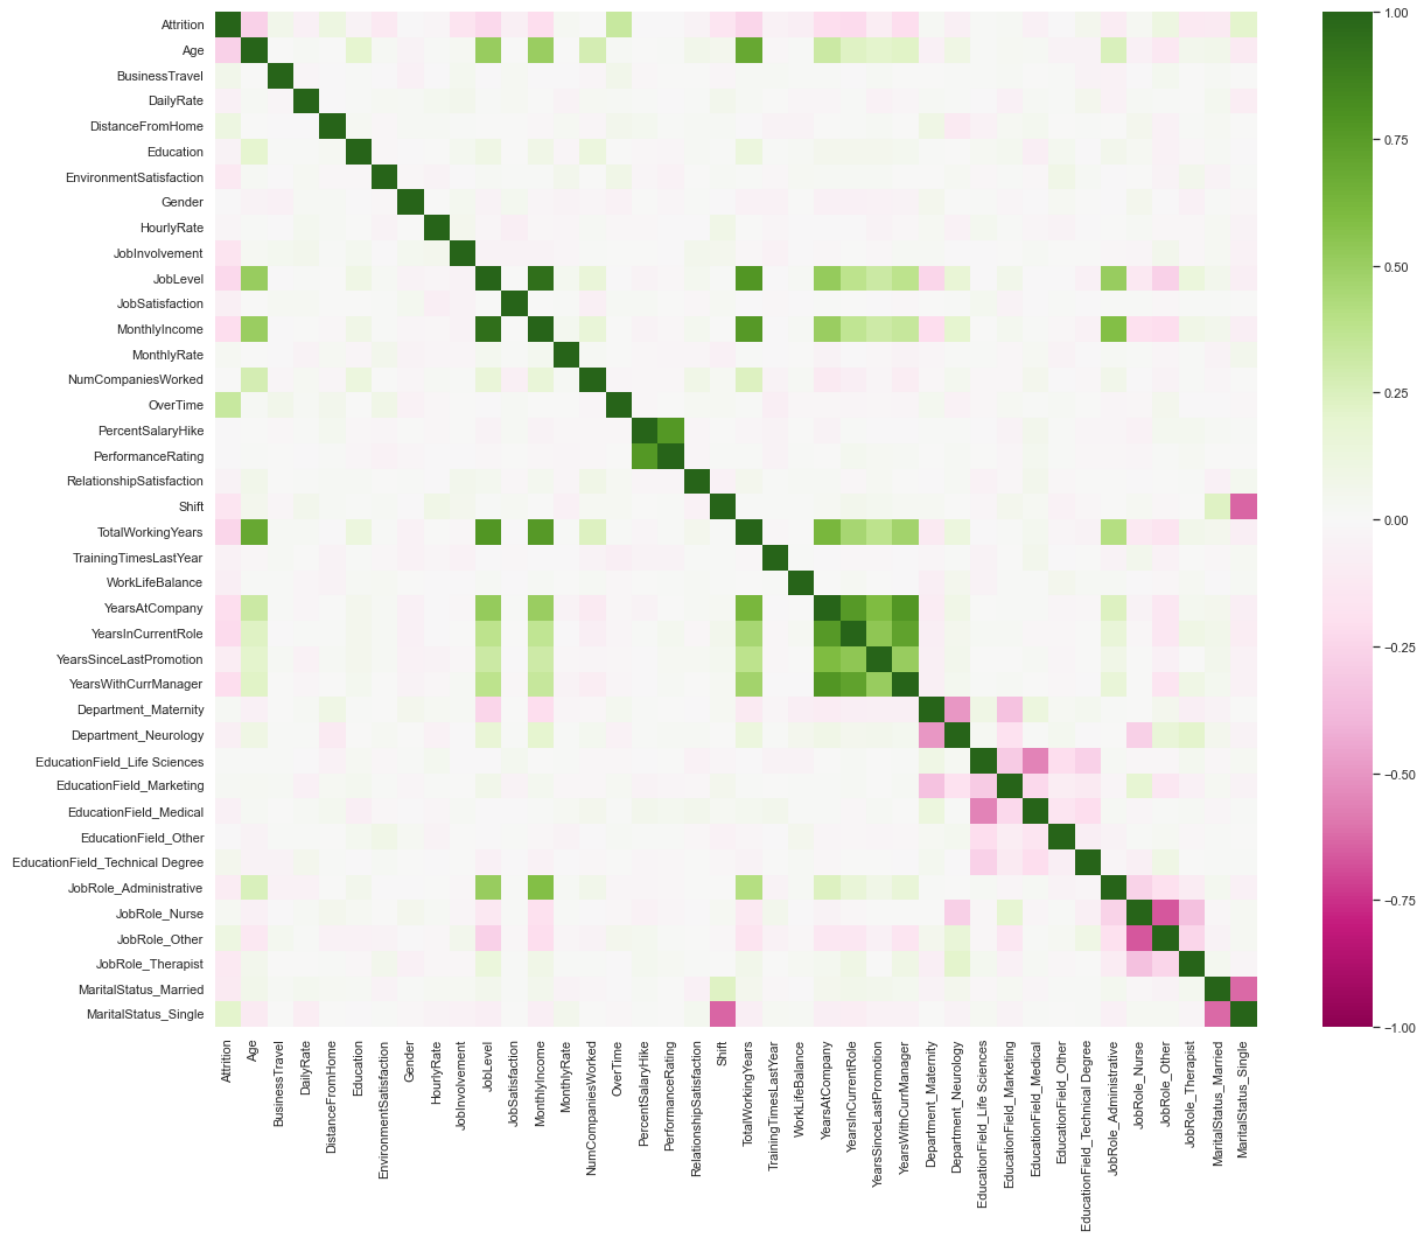
	Age	Attrition	BusinessTravel	DailyRate	DistanceFromHome	Education	Environment
count	1340.000000	1340.000000	1340.000000	1340.000000	1340.000000	1340.000000	1340.000000
mean	36.580597	0.118657	1.092537	799.197761	9.193284	2.924627	2.924627
std	9.013072	0.323505	0.534490	399.333256	8.141621	1.036088	1.036088
min	18.000000	0.000000	0.000000	102.000000	1.000000	1.000000	1.000000
25%	30.000000	0.000000	1.000000	465.000000	2.000000	2.000000	2.000000
50%	35.000000	0.000000	1.000000	796.000000	7.000000	3.000000	3.000000
75%	42.000000	0.000000	1.000000	1153.000000	14.000000	4.000000	4.000000
max	60.000000	1.000000	2.000000	1499.000000	29.000000	5.000000	5.000000

8 rows × 40 columns

In [103...

```
attrition = train_data.pop('Attrition')
train_data.insert(0, 'Attrition', attrition)

correlation_matrix = train_data.corr().round(2)
sns.heatmap(data = correlation_matrix, annot = False, cmap="PiYG", vmin=-1, vmax=1)
sns.set(rc = {'figure.figsize': (20,16)})
plt.show()
```

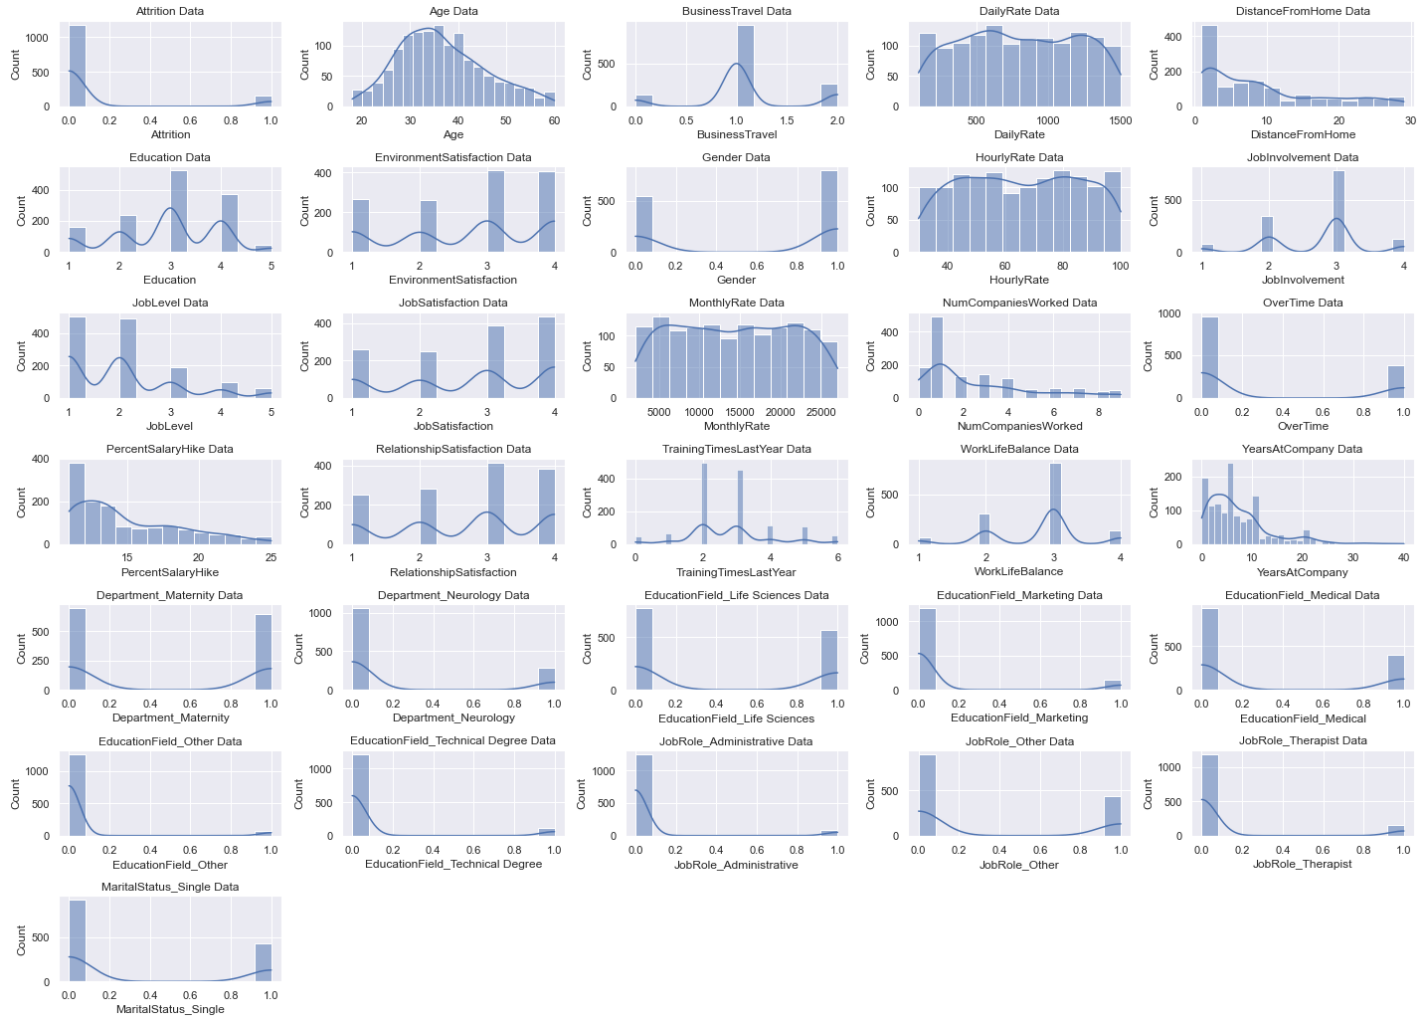


In [13]:

```
# drop data
train_data = train_data.drop(['MonthlyIncome', 'TotalWorkingYears', 'PerformanceRating', 'YearsSinceLastPromotion', 'YearsWithCurrManager', 'JobRole_Married', 'MaritalStatus_Married', 'Shift'], axis=1)
test_data = test_data.drop(['MonthlyIncome', 'TotalWorkingYears', 'PerformanceRating', 'YearsSinceLastPromotion', 'YearsWithCurrManager', 'JobRole_Married', 'MaritalStatus_Married', 'Shift'], axis=1)
```

In [14]:

```
# plt.figure(figsize=(12,21))
for i,col in enumerate(train_data.columns,1):
    plt.subplot(8,5,i)
    plt.title(f"{col} Data")
    sns.histplot(train_data[col],kde=True)
    plt.tight_layout()
    plt.plot()
```



In [57]:

```
# feature engineering
# log
def log_for_longtail(df):
    df['DistanceFromHome'] = np.log((1+ df['DistanceFromHome'] ))
    df['JobLevel'] = np.log((1+ df['JobLevel'] ))
    df['NumCompaniesWorked'] = np.log((1+ df['NumCompaniesWorked']))
    df['PercentSalaryHike'] = np.log((1+ df['PercentSalaryHike'] ))
    df['YearsAtCompany'] = np.log((1+ df['YearsAtCompany'] ))

    return df
```

In [104...]

```
# # feature engineering
# # log
# def log_for_longtail(df):
#     df['DistanceFromHome'] = np.log((1+ df['DistanceFromHome'] ))
#     df['JobLevel'] = np.log((1+ df['JobLevel'] ))
#     df['MonthlyIncome'] = np.log((1+ df['MonthlyIncome'] ))
#     df['NumCompaniesWorked'] = np.log((1+ df['NumCompaniesWorked']))
#     df['PercentSalaryHike'] = np.log((1+ df['PercentSalaryHike'] ))
#     df['TotalWorkingYears'] = np.log((1+ df['TotalWorkingYears'] ))
#     df['YearsInCurrentRole'] = np.log((1+ df['YearsInCurrentRole'] ))
#     df['YearsAtCompany'] = np.log((1+ df['YearsAtCompany'] ))
#     df['YearsSinceLastPromotion'] = np.log((1+df['YearsSinceLastPromotion'] ))
#     df['YearsWithCurrManager'] = np.log((1+ df['YearsWithCurrManager'] ))

#     return df
```

In [105...]

```
train_data = log_for_longtail(train_data)
test_data = log_for_longtail(test_data)
```

```
In [88]: print(len(train_data.columns))

40
```

feature selection

```
In [106... # split data into X and y (data and label)
X_nondrop = train_data.drop(columns=['Attrition'], axis = 1)
y = train_data['Attrition']
```

```
In [90]: X_nondrop_scaled = StandardScaler().fit_transform(X_nondrop)
```

```
In [48]: classes_count = y.value_counts()
classes_count
```

```
Out[48]: 0    1181
1     159
Name: Attrition, dtype: int64
```

```
In [28]: index = ['Logistic regression',
                  'Random Forest',
                  'Logistic regression with balanced class weights',
                  'Balanced random forest',
                  'Balanced bag of histogram gradient boosting']

scores = {"Accuracy": [], "Balanced accuracy": [], "f1": [], "f1_weighted":[]}

scoring = ["accuracy", "balanced_accuracy", "f1", "f1_weighted"]

clfs = [LogisticRegression(n_jobs=-1), RandomForestClassifier(n_jobs=-1),
        LogisticRegression(class_weight="balanced", n_jobs=-1),
        BalancedRandomForestClassifier(n_jobs=-1), BalancedBaggingClassifier(
            base_estimator=HistGradientBoostingClassifier(random_state=42),
            n_jobs=-1)]
```

```
In [29]: for i in clfs:
cv_result = cross_validate(i, X_nondrop_scaled, y, scoring=scoring)
scores["Accuracy"].append(cv_result["test_accuracy"].mean())
scores["Balanced accuracy"].append(cv_result["test_balanced_accuracy"].mean())
scores["f1"].append(cv_result["test_f1"].mean())
scores["f1_weighted"].append(cv_result["test_f1_weighted"].mean())

df_scores = pd.DataFrame(scores, index=index)
df_scores
```

Out[29]:

	Accuracy	Balanced accuracy	f1	f1_weighted
Logistic regression	0.935821	0.806290	0.694960	0.932069
Random Forest	0.914179	0.665722	0.479481	0.896930
Logistic regression with balanced class weights	0.869403	0.866068	0.610142	0.884530
Balanced random forest	0.857463	0.859197	0.589179	0.875258
Balanced bag of histogram gradient boosting	0.895522	0.869887	0.655377	0.904842



In [30]:

```
from sklearn.pipeline import Pipeline
from numpy import mean
from numpy import std
from sklearn.model_selection import StratifiedKFold
from sklearn.model_selection import cross_val_score
from sklearn.svm import SVC

from sklearn.feature_selection import RFECV

estimator_for_feature = [RandomForestClassifier(random_state=42)]

plt.figure(figsize=(8,6))
plt.figure()
plt.xlabel("Number of features selected")
plt.ylabel("Cross validation score (F1)")

for i in estimator_for_feature:
    rfecv = RFECV(
        estimator = i, scoring = "f1", step = 1, n_jobs=-1, verbose=0,
        cv = RepeatedStratifiedKFold(n_splits=5, n_repeats= 3, random_state=42),
        min_features_to_select=1)
    rfecv.fit(X_nondrop, y)

    print("%s Optimal number of features : %d, achieving F1: %.3f" %
          (i, rfecv.n_features_, mean(rfecv.grid_scores_[rfecv.n_features_])))


    # Plot number of features VS. cross-validation scores

    plt.plot(
        range(1, len(rfecv.grid_scores_) + 1),
        [mean(j) for j in rfecv.grid_scores_], label = str(i))

plt.legend()
plt.show()
```

RandomForestClassifier(random\_state=42) Optimal number of features : 12, achieving F1: 0.520

/Users/minghuizhu/opt/anaconda3/lib/python3.9/site-packages/sklearn/utils/deprecation.py:103: FutureWarning: The `grid\_scores\_` attribute is deprecated in version 1.0 in favor of `cv\_results\_` and will be removed in version 1.2.  
warnings.warn(msg, category=FutureWarning)  
<Figure size 576x432 with 0 Axes>



In [31]:

```
# define RFE
rfe_lg = RFE(estimator=RandomForestClassifier(random_state=42), n_features_to_select=12)

# fit RFE
rfe_lg.fit(X_nondrop, y)

# summarize all features
dict_for_feature_select_lg = dict()
dt = pd.DataFrame(columns=["Feature Name", "Selected"])

for i in range(X_nondrop.shape[1]):
    dict_for_feature_select_lg[X_nondrop.columns[i]] = rfe_lg.support_[i]
    new = [str(X_nondrop.columns[i]), str(rfe_lg.support_[i])]
    dt.loc[len(dt.index)] = new

print(dt)
```

	Feature Name	Selected
0	Age	True
1	BusinessTravel	False
2	DailyRate	True
3	DistanceFromHome	True
4	Education	False
5	EnvironmentSatisfaction	True
6	Gender	False
7	HourlyRate	True
8	JobInvolvement	True
9	JobLevel	True
10	JobSatisfaction	False

11	MonthlyRate	True
12	NumCompaniesWorked	True
13	OverTime	True
14	PercentSalaryHike	True
15	RelationshipSatisfaction	False
16	TrainingTimesLastYear	False
17	WorkLifeBalance	False
18	YearsAtCompany	True
19	Department_Maternity	False
20	Department_Neurology	False
21	EducationField_Life Sciences	False
22	EducationField_Marketing	False
23	EducationField_Medical	False
24	EducationField_Other	False
25	EducationField_Technical Degree	False
26	JobRole_Administrative	False
27	JobRole_Other	False
28	JobRole_Therapist	False
29	MaritalStatus_Single	False

In [32]:

```
dict_for_feature_select_list = []

for col in X_nondrop.columns:
    if dict_for_feature_select_lg[col] == False:
        dict_for_feature_select_list.append(col)

X_drop = X_nondrop.drop(dict_for_feature_select_list, axis=1)

print(X_drop.shape)
```

(1340, 12)

In [107]:

```
X = X_nondrop

# apply standard scaler onto the training data
scale = StandardScaler()
scale.fit(X)
X = scale.transform(X)
```

In [36]:

```
import json
from json import dumps
```

In [108]:

```
# source: Matt. Validation Curve Plot from GridSearchCV Results.
# Retrieved from https://matthewbilyeu.com/blog/2019-02-05/validation-curve-plot-from-gridsearchcv/
import numbers
def plot_grid_search_validation_curve(grid, param_to_vary,
                                     title='Validation Curve', ylim=None,
                                     xlim=None, log=None):
    """Plots train and cross-validation scores from a GridSearchCV instance's
    best params while varying one of those params."""

    df_cv_results = pd.DataFrame(grid.cv_results_)

    print(type(df_cv_results))
    train_scores_mean = df_cv_results['mean_train_score']
    valid_scores_mean = df_cv_results['mean_test_score']
    train_scores_std = df_cv_results['std_train_score']
    valid_scores_std = df_cv_results['std_test_score']

    param_cols = [c for c in df_cv_results.columns if c[:6] == 'param_']
    param_ranges = [grid.param_grid[p[6:]] for p in param_cols]
```

```

param_ranges_lengths = [len(pr) for pr in param_ranges]

train_scores_mean = np.array(train_scores_mean).reshape(*param_ranges_lengths)
valid_scores_mean = np.array(valid_scores_mean).reshape(*param_ranges_lengths)
train_scores_std = np.array(train_scores_std).reshape(*param_ranges_lengths)
valid_scores_std = np.array(valid_scores_std).reshape(*param_ranges_lengths)

param_to_vary_idx = param_cols.index('param_{}'.format(param_to_vary))

slices = []
for idx, param in enumerate(grid.best_params_):
    if (idx == param_to_vary_idx):
        slices.append(slice(None))
        continue
    best_param_val = grid.best_params_[param]
    idx_of_best_param = 0
    if isinstance(param_ranges[idx], np.ndarray):
        idx_of_best_param = param_ranges[idx].tolist().index(best_param_val)
    else:
        idx_of_best_param = param_ranges[idx].index(best_param_val)
    slices.append(idx_of_best_param)

train_scores_mean = train_scores_mean[tuple(slices)]
valid_scores_mean = valid_scores_mean[tuple(slices)]
train_scores_std = train_scores_std[tuple(slices)]
valid_scores_std = valid_scores_std[tuple(slices)]

plt.figure(figsize=(8,6))
plt.clf()

plt.title(title)
plt.xlabel(param_to_vary)
plt.ylabel('Score')

if (ylim is None):
    plt.ylim(0.0, 1.1)
else:
    plt.ylim(*ylim)

if (not (xlim is None)):
    plt.xlim(*xlim)

lw = 2

plot_fn = plt.plot
if log:
    plot_fn = plt.semilogx

param_range = param_ranges[param_to_vary_idx]
if (not isinstance(param_range[0], numbers.Number)):
    param_range = [str(x) for x in param_range]
plot_fn(param_range, train_scores_mean, label='Training score', color='r',
        lw=lw)
plt.fill_between(param_range, train_scores_mean - train_scores_std,
                train_scores_mean + train_scores_std, alpha=0.1,
                color='r', lw=lw)
plot_fn(param_range, valid_scores_mean, label='Cross-validation score',
        color='b', lw=lw)
plt.fill_between(param_range, valid_scores_mean - valid_scores_std,
                valid_scores_mean + valid_scores_std, alpha=0.1,
                color='b', lw=lw)

plt.legend(loc='lower right')

plt.show()

```

In [145...

```
# modeling
sample_weight = np.array([10 if i == 1 else 1 for i in y])
def hypertune_fit_eval(estimator, param_grid):

    grid_obj = GridSearchCV(estimator = estimator, param_grid = param_grid, return_train_score=True,
                             cv = StratifiedKFold(), verbose=0, n_jobs = -1, scoring = "f1")

    start = time.time()
    grid_obj.fit(X, y)
    end = time.time()

    print("\n The best estimator:\n", grid_obj.best_estimator_)
    print("\n The best score:\n", grid_obj.best_score_)
    print("\n The best parameters:\n", grid_obj.best_params_)
    print("\n Train Time:\n", end-start)

    random_id_for_kaggle = ''.join(random.choices(string.ascii_uppercase + string.digits,
                                                    k=10))
    timestamp = str(datetime.datetime.now())
    estimator = str(estimator)
    best_estimator_ = str(grid_obj.best_estimator_)
    train_time = str(end - start)
    best_f1 = str(grid_obj.best_score_)
    X_cols = dict_for_feature_select_list

    new_row = [random_id_for_kaggle, timestamp, estimator, best_estimator_,
               train_time, best_f1, X_cols]

    records = sa.open("records").worksheet("Sheet1")
    records_df = pd.DataFrame(records.get_all_records())

    nrow = len(records_df)

    for i in range(len(new_row)):
        records.update_cell(nrow+2, i+1, json.dumps(new_row[i]))

    return grid_obj
```

## Hypertune: LR

In [110...

```
# logistic regression
lr = LogisticRegression(random_state=42, class_weight="balanced")

lr_grid = {
    "solver": ['newton-cg', 'lbfgs', 'liblinear', 'sag', 'saga']
}

best_lr = hypertune_fit_eval(lr, lr_grid)
plot_grid_search_validation_curve(best_lr, 'solver', title='LogisticRegression')
```

```
/Users/minghuizhu/opt/anaconda3/lib/python3.9/site-packages/sklearn/linear_model/_sag.py:3
50: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
/Users/minghuizhu/opt/anaconda3/lib/python3.9/site-packages/sklearn/linear_model/_sag.py:3
50: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
/Users/minghuizhu/opt/anaconda3/lib/python3.9/site-packages/sklearn/linear_model/_sag.py:3
50: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
/Users/minghuizhu/opt/anaconda3/lib/python3.9/site-packages/sklearn/linear_model/_sag.py:3
50: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
/Users/minghuizhu/opt/anaconda3/lib/python3.9/site-packages/sklearn/linear_model/_sag.py:3
```

```

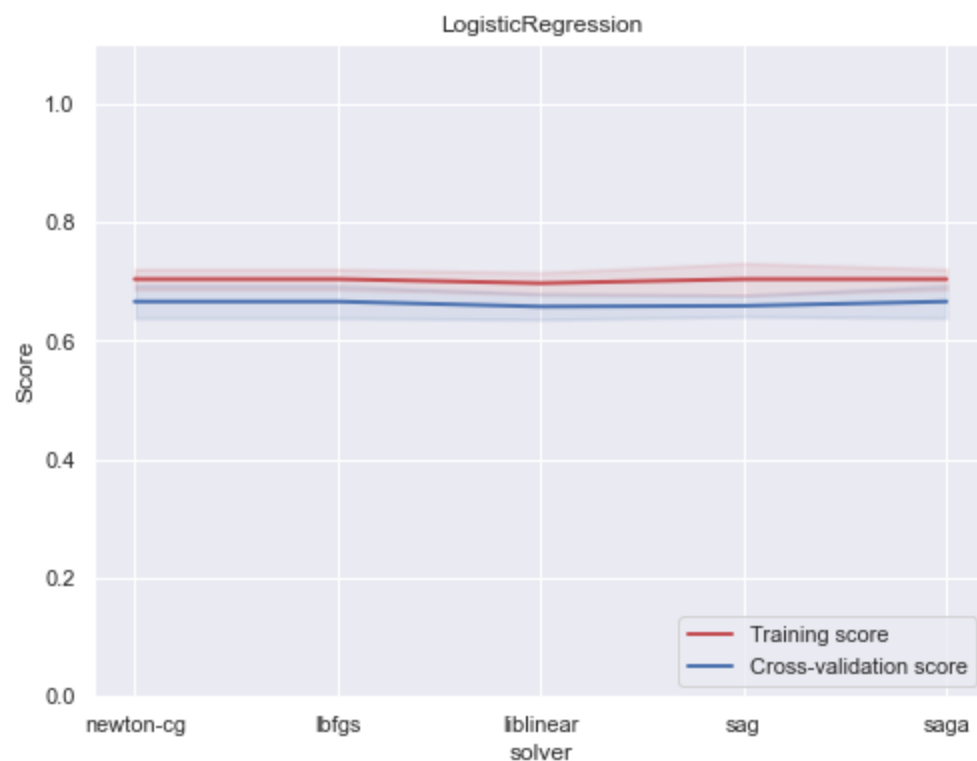
50: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
/Users/minghuizhu/opt/anaconda3/lib/python3.9/site-packages/sklearn/linear_model/_sag.py:3
50: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
/Users/minghuizhu/opt/anaconda3/lib/python3.9/site-packages/sklearn/linear_model/_sag.py:3
50: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
/Users/minghuizhu/opt/anaconda3/lib/python3.9/site-packages/sklearn/linear_model/_sag.py:3
50: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
/Users/minghuizhu/opt/anaconda3/lib/python3.9/site-packages/sklearn/linear_model/_sag.py:3
50: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
/Users/minghuizhu/opt/anaconda3/lib/python3.9/site-packages/sklearn/linear_model/_sag.py:3
50: ConvergenceWarning: The max_iter was reached which means the coef_ did not converge
warnings.warn(
The best estimator:
LogisticRegression(class_weight='balanced', random_state=42, solver='newton-cg')

The best score:
0.6660420007122451

The best parameters:
{'solver': 'newton-cg'}

Train Time:
0.17855310440063477
<class 'pandas.core.frame.DataFrame'>

```



In [111...

```

# logistic regression
lr = best_lr.best_estimator_

lr_grid = {
    "C": np.logspace(-4, 4, 50)
}

best_lr = hypertune_fit_eval(lr, lr_grid)
plot_grid_search_validation_curve(best_lr, 'C', title='LogisticRegression')

```


```
The best estimator:  
LogisticRegression(C=0.3906939937054613, class_weight='balanced',  
                    random_state=42, solver='newton-cg')
```

```
The best score:  
0.6728700486777743
```

```
The best parameters:  
{'C': 0.3906939937054613}
```

```
Train Time:  
0.7589290142059326
```

```
<class 'pandas.core.frame.DataFrame'>
```



In [160...

```
print(np.logspace(-4,4, 50))
```

```
[1.00000000e-04 1.45634848e-04 2.12095089e-04 3.08884360e-04  
4.49843267e-04 6.55128557e-04 9.54095476e-04 1.38949549e-03  
2.02358965e-03 2.94705170e-03 4.29193426e-03 6.25055193e-03  
9.10298178e-03 1.32571137e-02 1.93069773e-02 2.81176870e-02  
4.09491506e-02 5.96362332e-02 8.68511374e-02 1.26485522e-01  
1.84206997e-01 2.68269580e-01 3.90693994e-01 5.68986603e-01  
8.28642773e-01 1.20679264e+00 1.75751062e+00 2.55954792e+00  
3.72759372e+00 5.42867544e+00 7.90604321e+00 1.15139540e+01  
1.67683294e+01 2.44205309e+01 3.55648031e+01 5.17947468e+01  
7.54312006e+01 1.09854114e+02 1.59985872e+02 2.32995181e+02  
3.39322177e+02 4.94171336e+02 7.19685673e+02 1.04811313e+03  
1.52641797e+03 2.22299648e+03 3.23745754e+03 4.71486636e+03  
6.86648845e+03 1.00000000e+04]
```

In [113...

```
# logistic regression  
lr = best_lr.best_estimator_  
  
lr_grid = {  
    'max_iter' : [100, 1000, 2500, 5000]  
}  
  
best_lr = hypertune_fit_eval(lr, lr_grid)  
plot_grid_search_validation_curve(best_lr, 'max_iter', title='LogisticRegression')
```

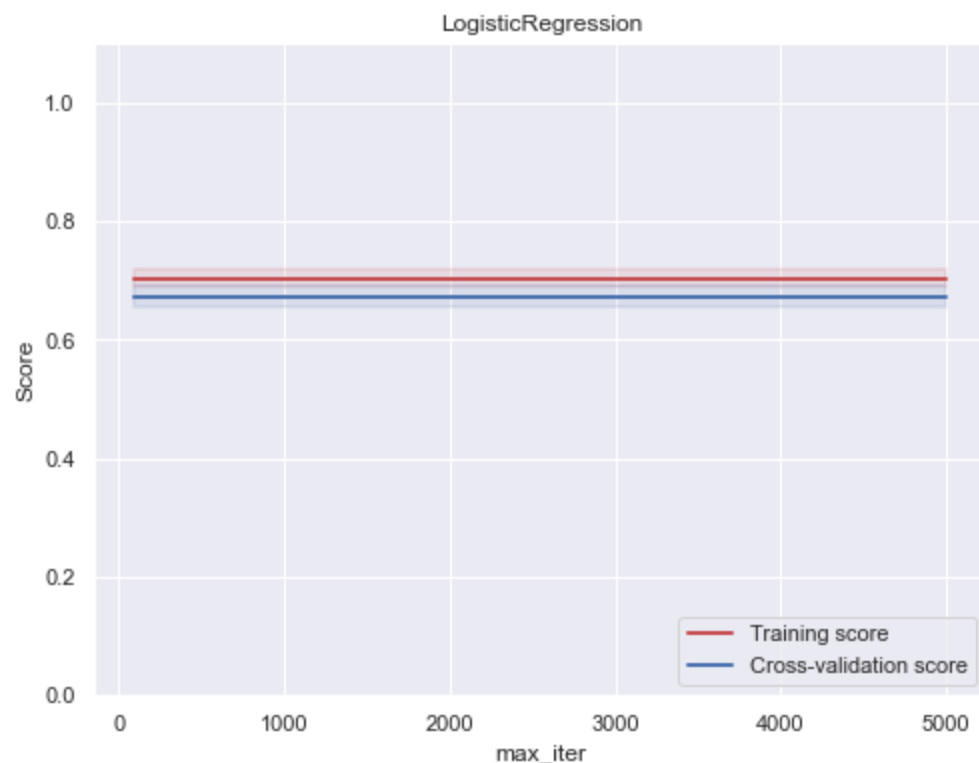
```
The best estimator:  
LogisticRegression(C=0.3906939937054613, class_weight='balanced',  
                    random_state=42, solver='newton-cg')
```

```
The best score:  
0.6728700486777743
```

```
The best parameters:  
{'max_iter': 100}
```

```
Train Time:  
0.11745715141296387
```

```
<class 'pandas.core.frame.DataFrame'>
```



In [ ]:

```
# logistic regression  
lr = best_lr.best_estimator_  
lr_grid = {  
    "penalty": ['l1', 'l2', 'elasticnet']  
}  
  
best_lr = hypertune_fit_eval(lr, lr_grid)  
plot_grid_search_validation_curve(best_lr, 'penalty', title='LogisticRegression')
```

In [97]:

```
# brf  
  
from imblearn.ensemble import BalancedRandomForestClassifier  
  
# random forest  
brf = BalancedRandomForestClassifier(n_jobs=-1, random_state=42)  
  
brf_grid = {'n_estimators': np.arange(100, 1500, 100)  
           }  
best_brf = hypertune_fit_eval(brf, brf_grid)  
plot_grid_search_validation_curve(best_brf, 'n_estimators', title='BalancedRandomForest')
```

```
The best estimator:  
BalancedRandomForestClassifier(n_estimators=400, n_jobs=-1, random_state=42)
```

```
The best score:
```




0.5938415713675168

The best parameters:  
{'n\_estimators': 400}

Train Time:  
30.14676809310913

<class 'pandas.core.frame.DataFrame'>



In [65]:

```
# brf

# random forest
brf = best_brf.best_estimator_

brf_grid = {
    'max_features': ["sqrt", "log2"],
}
best_brf = hypertune_fit_eval(brf, brf_grid)
plot_grid_search_validation_curve(best_brf, 'max_features', title='BalancedRandomForest')
```

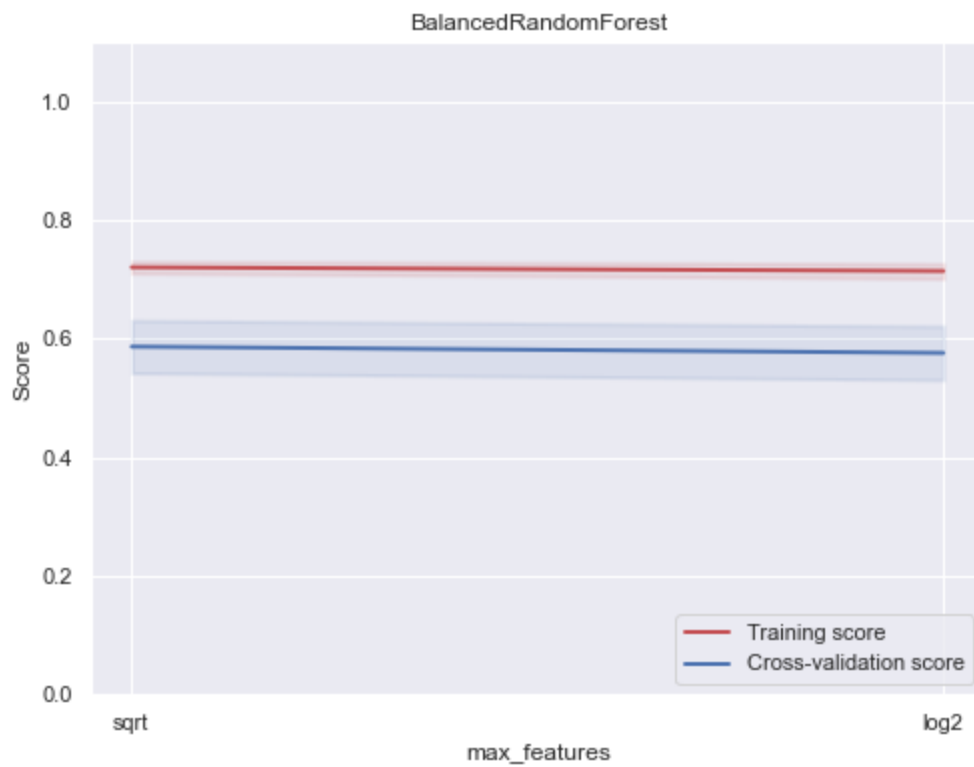
The best estimator:  
BalancedRandomForestClassifier(n\_estimators=500, n\_jobs=-1, random\_state=42)

The best score:  
0.5866386404694979

The best parameters:  
{'max\_features': 'sqrt'}

Train Time:  
4.356066942214966

<class 'pandas.core.frame.DataFrame'>



In [66]:

```
# brf

# random forest
brf = best_brf.best_estimator_

brf_grid = {
    'max_depth': np.arange(3,20,1)
}
best_brf = hypertune_fit_eval(brf, brf_grid)
plot_grid_search_validation_curve(best_brf, 'max_depth', title='BalancedRandomForest')
```

The best estimator:

```
BalancedRandomForestClassifier(max_depth=9, n_estimators=500, n_jobs=-1,
                                random_state=42)
```

The best score:

```
0.5900157351880037
```


The best parameters:

```
{'max_depth': 9}
```

Train Time:

```
24.133669137954712
```

```
<class 'pandas.core.frame.DataFrame'>
```



In [67]:

```
# brf

# random forest
brf = best_brf.best_estimator_

brf_grid = {
    'min_samples_split': np.arange(1,5,1),
}

best_brf = hypertune_fit_eval(brf, brf_grid)
plot_grid_search_validation_curve(best_brf, 'min_samples_split', title='BalancedRandomFore
```

/Users/minghuizhu/opt/anaconda3/lib/python3.9/site-packages/sklearn/model\_selection/\_validation.py:378: FitFailedWarning:  
5 fits failed out of a total of 20.  
The score on these train-test partitions for these parameters will be set to nan.  
If these failures are not expected, you can try to debug them by setting error\_score='raise'.

Below are more details about the failures:

```
-----
5 fits failed with the following error:
Traceback (most recent call last):
  File "/Users/minghuizhu/opt/anaconda3/lib/python3.9/site-packages/sklearn/model_selection/_validation.py", line 686, in _fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "/Users/minghuizhu/opt/anaconda3/lib/python3.9/site-packages/imblearn/ensemble/_forest.py", line 547, in fit
    samplers_trees = Parallel(
  File "/Users/minghuizhu/opt/anaconda3/lib/python3.9/site-packages/joblib/parallel.py", line 1056, in __call__
    self.retrieve()
  File "/Users/minghuizhu/opt/anaconda3/lib/python3.9/site-packages/joblib/parallel.py", line 935, in retrieve
    self._output.extend(job.get(timeout=self.timeout))
  File "/Users/minghuizhu/opt/anaconda3/lib/python3.9/multiprocessing/pool.py", line 771, in get
    raise self._value
  File "/Users/minghuizhu/opt/anaconda3/lib/python3.9/multiprocessing/pool.py", line 125,
```

```

in worker
    result = (True, func(*args, **kwargs))
    File "/Users/minghuizhu/opt/anaconda3/lib/python3.9/site-packages/joblib/_parallel_backe
nds.py", line 595, in __call__
        return self.func(*args, **kwargs)
    File "/Users/minghuizhu/opt/anaconda3/lib/python3.9/site-packages/joblib/parallel.py", 1
ine 262, in __call__
        return [func(*args, **kwargs)
    File "/Users/minghuizhu/opt/anaconda3/lib/python3.9/site-packages/joblib/parallel.py", 1
ine 262, in <listcomp>
        return [func(*args, **kwargs)
    File "/Users/minghuizhu/opt/anaconda3/lib/python3.9/site-packages/imblearn/ensemble/_for
est.py", line 61, in _local_parallel_build_trees
        tree = _parallel_build_trees(
    File "/Users/minghuizhu/opt/anaconda3/lib/python3.9/site-packages/sklearn/ensemble/_fore
st.py", line 189, in _parallel_build_trees
        tree.fit(X, y, sample_weight=curr_sample_weight, check_input=False)
    File "/Users/minghuizhu/opt/anaconda3/lib/python3.9/site-packages/sklearn/tree/_classes.
py", line 969, in fit
        super().fit(
    File "/Users/minghuizhu/opt/anaconda3/lib/python3.9/site-packages/sklearn/tree/_classes.
py", line 265, in fit
        check_scalar(
    File "/Users/minghuizhu/opt/anaconda3/lib/python3.9/site-packages/sklearn/utils/validati
on.py", line 1480, in check_scalar
        raise ValueError(
ValueError: min_samples_split == 1, must be >= 2.

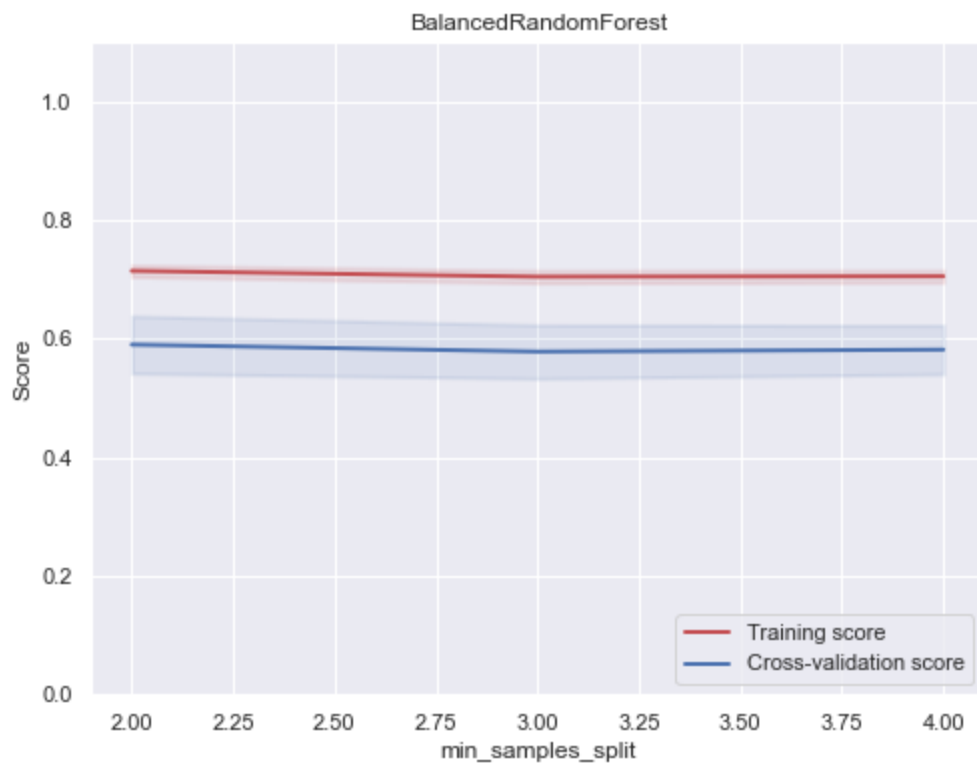
    warnings.warn(some_fits_failed_message, FitFailedWarning)
/Users/minghuizhu/opt/anaconda3/lib/python3.9/site-packages/sklearn/model_selection/_searc
h.py:953: UserWarning: One or more of the test scores are non-finite: [          nan 0.590015
74 0.57822054 0.58156182]
    warnings.warn(
/Users/minghuizhu/opt/anaconda3/lib/python3.9/site-packages/sklearn/model_selection/_searc
h.py:953: UserWarning: One or more of the train scores are non-finite: [          nan 0.71436
58 0.70451168 0.70526346]
    warnings.warn(
The best estimator:
BalancedRandomForestClassifier(max_depth=9, n_estimators=500, n_jobs=-1,
                                random_state=42)

The best score:
0.5900157351880037

The best parameters:
{'min_samples_split': 2}

Train Time:
4.845700025558472
<class 'pandas.core.frame.DataFrame'>

```



In [68]:

```
# brf

# random forest
brf = best_brf.best_estimator_

brf_grid = {
    'min_samples_leaf': np.arange(1, 5, 1)
}
best_brf = hypertune_fit_eval(brf, brf_grid)
plot_grid_search_validation_curve(best_brf, 'min_samples_leaf', title='BalancedRandomFore
```

The best estimator:

BalancedRandomForestClassifier(max\_depth=9, n\_estimators=500, n\_jobs=-1,  
random\_state=42)

The best score:

0.5900157351880037

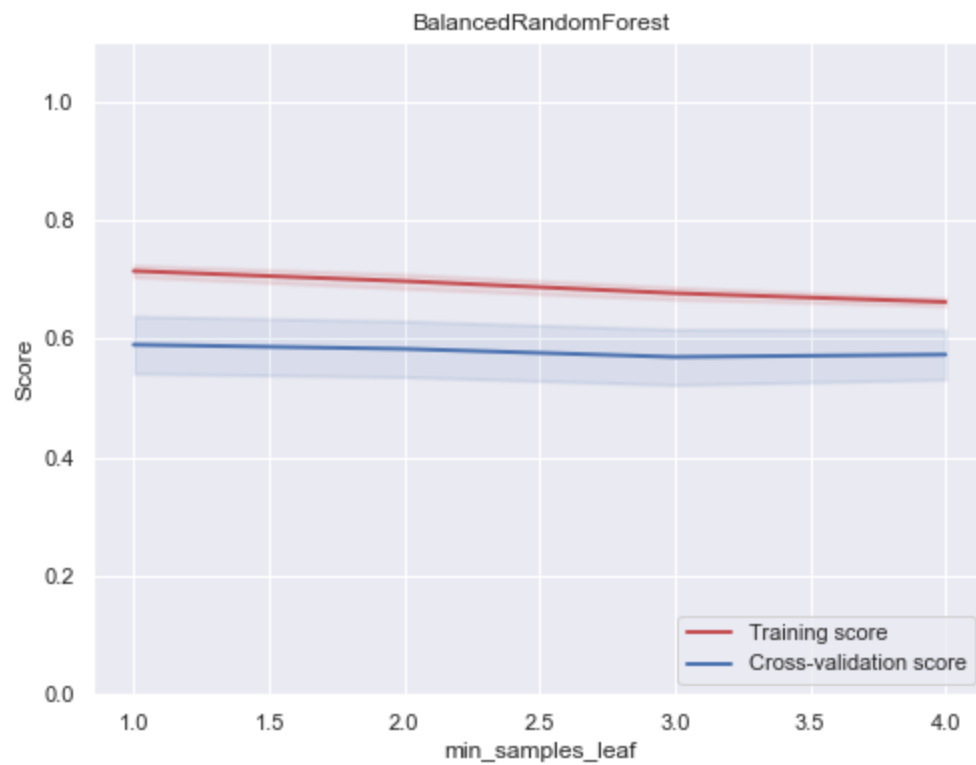
The best parameters:

{'min\_samples\_leaf': 1}

Train Time:

6.5085930824279785

<class 'pandas.core.frame.DataFrame'>



In [70]:

```
# brf

# random forest
brf = best_brf.best_estimator_

brf_grid = {
    'bootstrap': [True, False]
}
best_brf = hypertune_fit_eval(brf, brf_grid)
plot_grid_search_validation_curve(best_brf, 'bootstrap', title='BalancedRandomForest')
```

The best estimator:

```
BalancedRandomForestClassifier(bootstrap=False, max_depth=11, n_estimators=900,
                                n_jobs=-1, random_state=42)
```

The best score:

```
0.6028438503260122
```


The best parameters:

```
{'bootstrap': False}
```

Train Time:

```
6.63654088973999
```

```
<class 'pandas.core.frame.DataFrame'>
```



In [240..

```
# gbc
gbc = GradientBoostingClassifier(random_state = 42)
gbc_grid = {
    "n_estimators": np.arange(50, 500, 50),
    "max_depth": [1, 3, 5, 7, 9, 11, 13],
    "learning_rate": np.arange(0.01, 0.3, 0.01),
    # 'colsample_bytree': np.arange(0.6, 1.0, 0.1),
    # 'colsample_bylevel': np.arange(0.6, 1.0, 0.1),
    'subsample': np.arange(0.6, 1.0, 0.1),
}
best_gbc = hypertune_fit_eval(gbc, gbc_grid)
```

```
-----
KeyboardInterrupt                                Traceback (most recent call last)
/var/folders/xq/07znzf8d6vd8p3m419_d9rv00000gn/T/ipykernel_1427/3124940627.py in <module>
      9     'subsample': np.arange(0.6, 1.0, 0.1),
     10 }
--> 11 best_gbc = hypertune_fit_eval(gbc, gbc_grid)

/var/folders/xq/07znzf8d6vd8p3m419_d9rv00000gn/T/ipykernel_1427/1032408374.py in hypertune
_fit_eval(estimator, param_grid)
      5                                     cv = StratifiedKFold(), verbose=0, n_jobs = -1, scor
ing = "f1", refit = "f1")
      6     start = time.time()
--> 7     grid_obj.fit(X, y)
      8     end = time.time()
      9

~/opt/anaconda3/lib/python3.9/site-packages/sklearn/model_selection/_search.py in fit(self, X, y, groups, **fit_params)
     873         return results
     874
--> 875         self._run_search(evaluate_candidates)
     876
     877         # multimetric is determined here because in the case of a callable

~/opt/anaconda3/lib/python3.9/site-packages/sklearn/model_selection/_search.py in _run_sea
rch(self, evaluate_candidates)
    1377     def _run_search(self, evaluate_candidates):
```

```

1378 """Search all candidates in param_grid"""
-> 1379     evaluate_candidates(ParameterGrid(self.param_grid))
1380
1381

~/opt/anaconda3/lib/python3.9/site-packages/sklearn/model_selection/_search.py in evaluate
_candidates(candidate_params, cv, more_results)
    820         )
    821
--> 822         out = parallel(
    823             delayed(_fit_and_score)(
    824                 clone(base_estimator),

~/opt/anaconda3/lib/python3.9/site-packages/joblib/parallel.py in __call__(self, iterable)
    1054
    1055         with self._backend.retrieval_context():
-> 1056             self.retrieve()
    1057         # Make sure that we get a last message telling us we are done
    1058         elapsed_time = time.time() - self._start_time

~/opt/anaconda3/lib/python3.9/site-packages/joblib/parallel.py in retrieve(self)
    933         try:
    934             if getattr(self._backend, 'supports_timeout', False):
--> 935                 self._output.extend(job.get(timeout=self.timeout))
    936             else:
    937                 self._output.extend(job.get())

~/opt/anaconda3/lib/python3.9/site-packages/joblib/_parallel_backends.py in wrap_future_re
sult(future, timeout)
    540         AsyncResults.get from multiprocessing."""
    541         try:
--> 542             return future.result(timeout=timeout)
    543         except CfTimeoutError as e:
    544             raise TimeoutError from e

~/opt/anaconda3/lib/python3.9/concurrent/futures/_base.py in result(self, timeout)
    438         return self.__get_result()
    439
--> 440         self._condition.wait(timeout)
    441
    442         if self._state in [CANCELLED, CANCELLED_AND_NOTIFIED]:

~/opt/anaconda3/lib/python3.9/threading.py in wait(self, timeout)
    310         try: # restore state no matter what (e.g., KeyboardInterrupt)
    311             if timeout is None:
--> 312                 waiter.acquire()
    313                 gotit = True
    314             else:

```

KeyboardInterrupt:

In [ ]:

```

# xgb
xgboost = XGBClassifier()
# Define the search space
param_grid = {
    # Learning rate shrinks the weights to make the boosting process more conservative
    # "learning_rate": [0.0001, 0.001, 0.01, 0.1, 1] ,
    # Maximum depth of the tree, increasing it increases the model complexity.
    # "max_depth": range(3, 21, 3),
    # Gamma specifies the minimum loss reduction required to make a split.
    "gamma": [i/10.0 for i in range(0, 5)],
    # Percentage of columns to be randomly sampled for each tree.
    "colsample_bytree": [i/10.0 for i in range(3, 10)],
    # reg_alpha provides l1 regularization to the weight, higher values result in more co
    "reg_alpha": [1e-5, 1e-2, 0.1, 1, 10, 100],

```



```

# reg_lambda provides 12 regularization to the weight, higher values result in more co
"reg_lambda": [1e-5, 1e-2, 0.1, 1, 10, 100],

}

xgb_model_tuned = hypertune_fit_eval(xgboost, param_grid)

```

```

In [71]: bbc = BalancedBaggingClassifier(base_estimator=HistGradientBoostingClassifier(random_state=42),
                                         n_jobs=-1)

bbc_grid = {
    "n_estimators": np.arange(50, 500, 50),
    # "max_depth": [1, 3, 5, 7, 9, 11, 13],
    # "learning_rate": np.arange(0.01, 0.3, 0.01),
    # 'colsample_bytree': np.arange(0.6, 1.0, 0.1),
    # 'colsample_bylevel': np.arange(0.6, 1.0, 0.1),
    # 'subsample': np.arange(0.6, 1.0, 0.1),
}

best_bbc = hypertune_fit_eval(bbc, bbc_grid)
plot_grid_search_validation_curve(best_bbc, 'n_estimators', title='BalancedBaggingClassifier')

```

The best estimator:

```
BalancedBaggingClassifier(base_estimator=HistGradientBoostingClassifier(random_state=42),
                          n_estimators=400, n_jobs=-1, random_state=42)
```

The best score:

```
0.6757341052521776
```


The best parameters:

```
{'n_estimators': 400}
```

Train Time:

```
311.64144587516785
```

```
<class 'pandas.core.frame.DataFrame'>
```



```

In [76]: bbc = best_bbc.best_estimator_

```

```
bbc_grid = {
    'base_estimator__max_depth': [1,3,5,7,9,11,13]
    # "learning_rate":np.arange(0.01, 0.3, 0.01),
    # 'colsample_bytree': np.arange(0.6, 1.0, 0.1),
    # 'colsample_bylevel': np.arange(0.6, 1.0, 0.1),
    # 'subsample': np.arange(0.6, 1.0, 0.1),
}
```

```
best_bbc = hypertune_fit_eval(bbc, bbc_grid)
plot_grid_search_validation_curve(best_bbc, 'base_estimator__max_depth', title='BalancedBaggingClassifier')
```

The best estimator:

```
BalancedBaggingClassifier(base_estimator=HistGradientBoostingClassifier(max_depth=9,
                                                                    random_state=42),
                          n_estimators=400, n_jobs=-1, random_state=42)
```

The best score:

```
0.6757341052521776
```

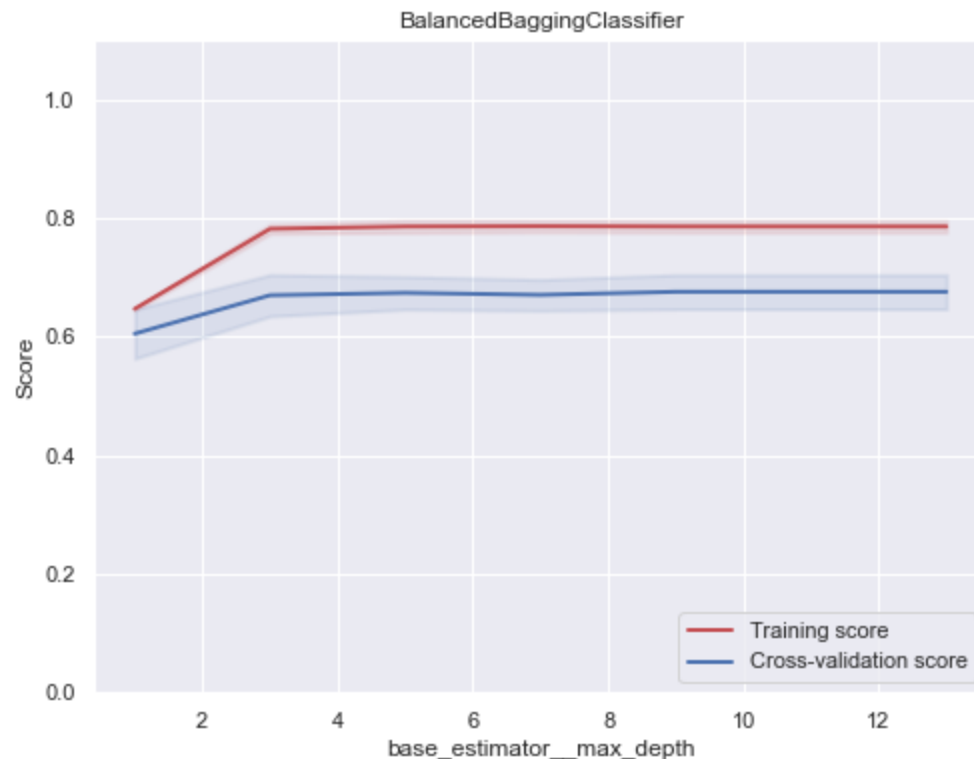
The best parameters:

```
{'base_estimator__max_depth': 9}
```

Train Time:

```
353.4232859611511
```

```
<class 'pandas.core.frame.DataFrame'>
```



In [116...]

```
bbc = best_bbc.best_estimator_

bbc_grid = {
    "base_estimator__learning_rate":np.arange(0.01, 0.3, 0.01),
}

best_bbc = hypertune_fit_eval(bbc, bbc_grid)
plot_grid_search_validation_curve(best_bbc, 'base_estimator__learning_rate', title='BalancedBaggingClassifier')
```

The best estimator:

```
BalancedBaggingClassifier(base_estimator=HistGradientBoostingClassifier(learning_rate=0.2,
                                                                    max_depth=9,
                                                                    random_state=42),
                          n_estimators=400, n_jobs=-1, random_state=42)
```

```
n_estimators=400, n_jobs=-1, random_state=42)
```

The best score:

0.6790538188912173

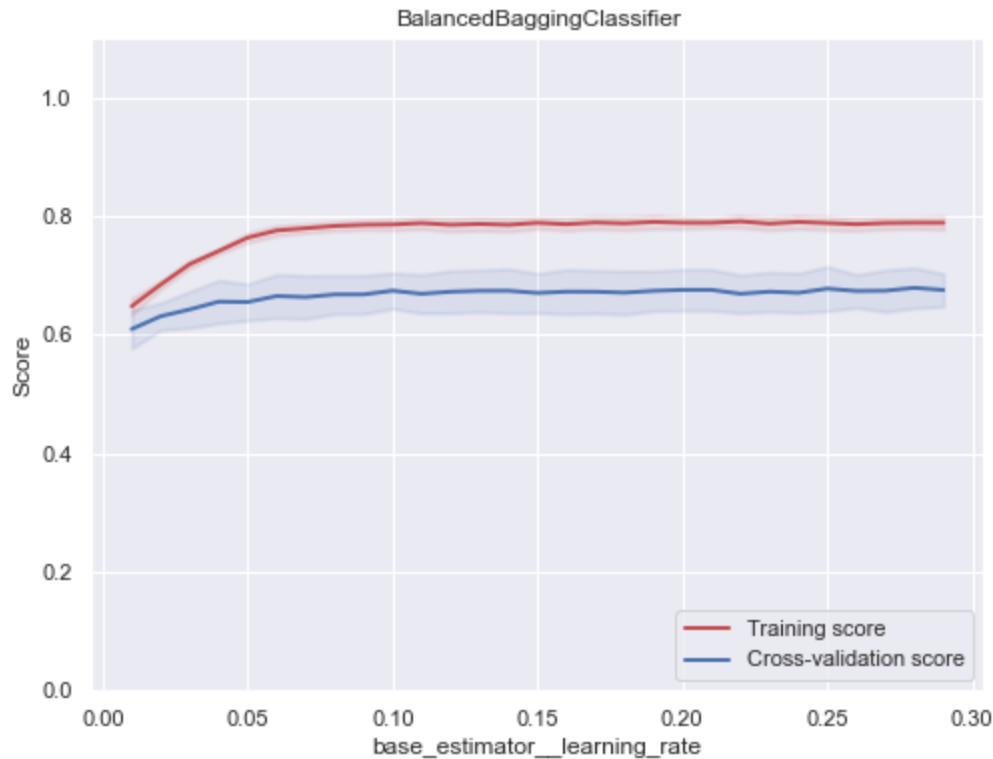
The best parameters:

```
{'base_estimator__learning_rate': 0.28}
```

Train Time:

1648.3858020305634

<class 'pandas.core.frame.DataFrame'>



In [334...

```
# adaboost
ada = AdaBoostClassifier(random_state = 42)
ada_grid = {
    "n_estimators": [int(i) for i in np.linspace(1, 100, num = 20)],
    "learning_rate": [0.01, 0.05, 0.1, 0.2, 0.4, 0.5, 1],
    "algorithm": ['SAMME', 'SAMME.R']
}
best_ada = hypertune_fit_eval(ada, ada_grid)
```

The best estimator:

AdaBoostClassifier(learning\_rate=0.5, n\_estimators=84, random\_state=42)

The best score:

0.7151626111324687

The best parameters:

```
{'algorithm': 'SAMME.R', 'learning_rate': 0.5, 'n_estimators': 84}
```

Train Time:

20.870769739151

In [ ]:

```
# easy ensemble
eec = EasyEnsembleClassifier(random_state=42, n_jobs=-1, verbose = 1, base_estimator = GradientDescentClassifier)
eec_grid = {
    "n_estimators": [int(i) for i in np.linspace(1, 100, num = 20)],
    # "learning_rate": [0.01, 0.05, 0.1, 0.2, 0.4, 0.5, 1, 5, 10, 15, 20, 30, 40, 50, 60, 70, 80, 90, 100],
    "sampling_strategy": ['float', 'str', 'dict', 'callable', 'auto']
}
```

```

}
best_eec = hypertune_fit_eval(eec, eec_grid)

```

In [138...

```

# svm
svm = SVC(random_state=42, kernel="sigmoid")
svm_grid = {
    'C': np.arange(5,30,5)
    # "gamma": [0.0001, 0.001, 0.01, 0.1, 1, 10, 100],
    # param_grid = {'C': [0.1,1, 10, 100], 'gamma': [1,0.1,0.01,0.001], 'kernel': ['rbf',
    # ]

best_svm = hypertune_fit_eval(svm, svm_grid)
plot_grid_search_validation_curve(best_svm, 'C', title='SVM', ylim=[0.4,0.8])

```


The best estimator:  
SVC(C=5, kernel='sigmoid', random\_state=42)

The best score:  
0.6556402737047898

The best parameters:  
{'C': 5}

Train Time:  
0.20191526412963867

<class 'pandas.core.frame.DataFrame'>



In [159...

```

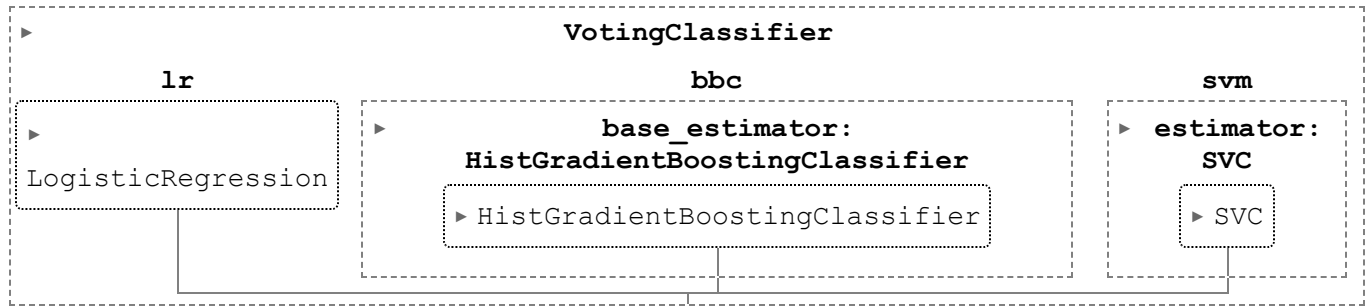
# try using voting methods
from sklearn.ensemble import VotingClassifier
from sklearn.ensemble import StackingClassifier
vo = [
    ("lr", best_lr.best_estimator_),
    ("bbc", best_bbc.best_estimator_),
    ("svm", best_svm_smote)]

clf = VotingClassifier(estimators=vo)

```

```
clf.fit(X,y)
```

Out [159...



In [141...

```
# svm
svm = SVC(random_state=42, kernel="rbf")
svm_grid = {
#     "gamma": [0.0001, 0.001, 0.01, 0.1, 1, 10, 100],
    'C': np.arange(1,10,1)
#     param_grid = {'C': [0.1,1, 10, 100], 'gamma': [1,0.1,0.01,0.001], 'kernel': ['rbf'],
}

best_svm_smote = hypertune_fit_eval(svm, svm_grid)
plot_grid_search_validation_curve(best_svm_smote, 'C', title='SVM smote')
```

The best estimator:  
SVC(C=4, random\_state=42)

The best score:  
0.6369620575502928

The best parameters:  
{'C': 4}

Train Time:  
0.8349130153656006

<class 'pandas.core.frame.DataFrame'>



In [155...

```
# svm
```

```

svm = SVC(random_state=42, kernel="rbf", C=4)
svm_grid = {
    "gamma": [0.0001, 0.001, 0.01, 0.1, 1, 10, 100]
}

best_svm_smote = hypertune_fit_eval(svm, svm_grid)
plot_grid_search_validation_curve(best_svm_smote, 'gamma', title='SVM smote')

```

The best estimator:

SVC(C=4, gamma=0.01, random\_state=42)

The best score:

0.6770558608058608

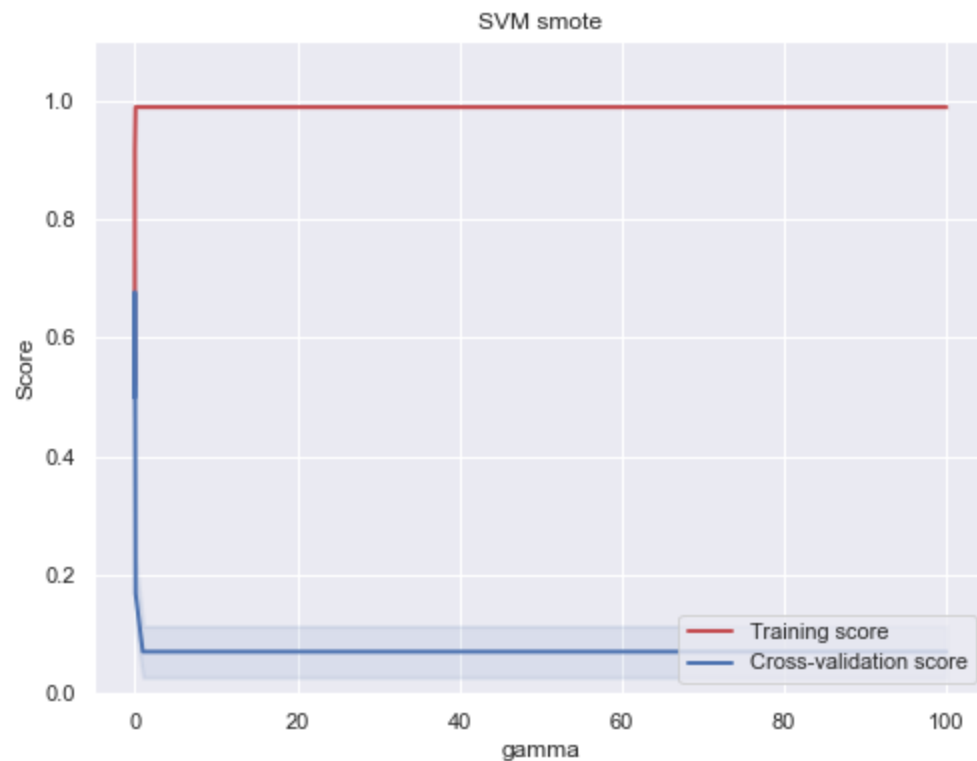
The best parameters:

{'gamma': 0.01}

Train Time:

1.5110981464385986

<class 'pandas.core.frame.DataFrame'>



In [74]:

```

def create_submission(grid, output='my_submission.csv'):
    clf = grid.best_estimator_
    # define the submission file
    submission_result = pd.DataFrame()
    # load test_data
    X_submission = test_data
    # X_submission = test_data.drop(dict_for_feature_select_list, axis=1)
    # apply standardscaler fitted earlier
    X_submission = scale.transform(X_submission)
    y_submission = clf.predict(X_submission)
    # record the prediction
    submission_result['Id'] = sample_submission_data['Id']
    submission_result['Predicted'] = y_submission
    # save submission
    submission_result.to_csv(output, index = False)

```

In [454...]

```

def create_submission_nondrop(grid, output='my_submission.csv'):
    clf = grid.best_estimator_

```

```
# define the submission file
submission_result = pd.DataFrame()
# load test_data
# X_submission = test_data
X_submission = test_data
# apply standard scaler fitted earlier
X_submission = scale.transform(X_submission)
y_submission = clf.predict(X_submission)
# record the prediction
submission_result['Id'] = sample_submission_data['Id']
submission_result['Predicted'] = y_submission
# save submission
submission_result.to_csv(output, index = False)
```

```
In [ ]: # rf
create_submission(rf_model_tuned, output='XAJI0Y6DPB.csv')
```

```
In [ ]: # xgb
create_submission(xgb_model_tuned, output='HSAHXTHV3A.csv')
```

```
In [115... # lg
create_submission(best_lr, output='best_lr.csv')
```

```
In [539... # lg nondrop
create_submission(best_lr, output='XXXXXXXX.csv')
```

```
In [ ]: # gbc
create_submission(best_model_gbc, output='H75LXO6QJI.csv')
```

```
In [335... # ada
create_submission(best_ada, output='31Y2R2FLJV.csv')
```

```
In [ ]: # eec
create_submission(best_eec, output='0Y9DOM5IGQ.csv')
```

```
In [110... # rus
create_submission(best_rus, output='RFZ7JM85VS.csv')
```

```
In [276... # svm
create_submission(best_svm, output='DEQU0ZEYV9.csv')
```

```
In [332... # svm smote
create_submission(best_svm_smote, output='WMN25BIKYZ.csv')
```

```
In [77]: # brf
create_submission(best_brf, output='best_brf.csv')
```

```
In [117... # bbc
create_submission(best_bbc, output='best_bbc.csv')
```

In [158...

```
def create_submission_voting(clf, output='my_submission.csv'):
    clf = clf
    # define the submission file
    submission_result = pd.DataFrame()
    # load test_data
    X_submission = test_data
    X_submission = scale.transform(X_submission)
    y_submission = clf.predict(X_submission)
    # record the prediction
    submission_result['Id'] = sample_submission_data['Id']
    submission_result['Predicted'] = y_submission
    # save submission
    submission_result.to_csv(output, index = False)

# bbc
create_submission_voting(clf, output='best_vog.csv')
```